

A deeper Look at Logging

Europython 2015, Bilbao

Stefan Baerisch

```
2015-07-13 15:06:16,145 INFO root
    : Doing division iteration 9995
2015-07-13 15:06:16,145 DEBUG root
    : Division 1 -31.591579612912746/-30.29
55474548343
[17:41:51] ~% allfiles>0021_traveld
ences_workshops>2015-07-13
iteration
```

Agenda

- Why Logging
- How does Loging work for you?
 - Components and structure
 - Getting the config right
 - Some special cases
- Optional Content

A simple program, without any logging

```
from datetime import datetime

def my_division_p(dividend, divisor):
    try:
        print("Debug, Division : {} / {}".format(dividend, divisor))
        result = dividend / divisor
    return result
except (ZeroDivisionError, TypeError):
    print("Error, Division Failed")
    return None

def division_task_handler_p(task):
    print("Handling division task, {} items".format(len(task)))
    result = []
    for i, task in enumerate(task):
        print("Doing devision iteration {} on {}".format(i, datetime.now()))
        dividend, divisor = task
        result.append(my_division_p(dividend, divisor))
    return result
```

Let us have look at the results

```
task = [(3,4),(5,1.4),(2,0),(3,5),("10",1)]  
division_task_handler_p(task)
```

```
Handling division task,5 items  
Doing devision iteration 0 on 2015  
Debug, Division : 3/4  
Doing devision iteration 1 on 2015  
Debug, Division : 5/1.4  
Doing devision iteration 2 on 2015  
Debug, Division : 2/0  
Error, Division Failed  
Doing devision iteration 3 on 2015  
Debug, Division : 3/5  
Doing devision iteration 4 on 2015  
Debug, Division : 10/1  
Error, Division Failed  
  
[0.75, 3.5714285714285716, None, 0.6, None]
```

The problems with using print()

- We don't have a way to centrally select the types of messages we are interested
 - By sources
 - By severity
- We have to add all information (timestamps, etc...) by ourselves
 - This is extra work
 - And all our messages look slightly different
- We have only limited control where our message end up

Using the the logging module for comparison

```
import log1; logging = log1.get_clean_logging()
logging.basicConfig(level=logging.DEBUG)
log = logging.getLogger()

def my_division(dividend, divisor):
    try:
        log.debug("Division : %s/%s", dividend, divisor)
        result = dividend / divisor
        return result
    except (ZeroDivisionError, TypeError):
        log.exception("Error, Division Failed")
        return None

def division_task_handler(task):
    log.info("Handling division task,%s items",len(task))
    result = []
    for i, task in enumerate(task):
        log.info("Doing devision iteration %s",i)
        dividend, divisor = task
        result.append(my_division(dividend,divisor))
    return result
```

The call and the Log Mesages

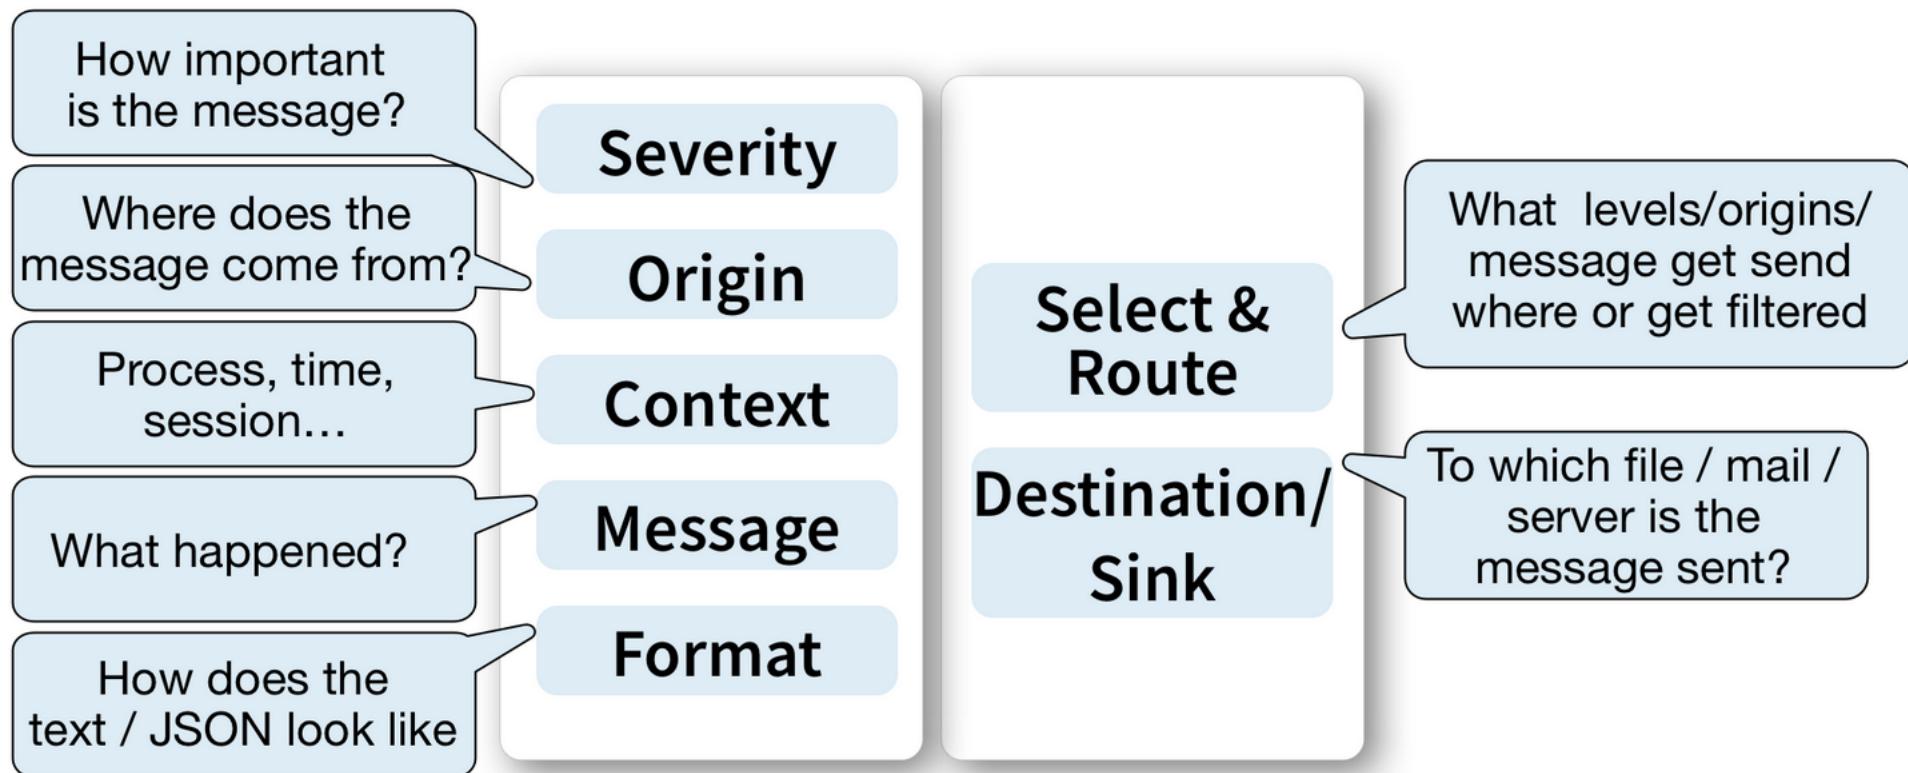
```
task = [(3,4),(2,0),(3,5),("10",1)]
division_task_handler(task)
```

```
INFO:root:Handling division task,4 items
INFO:root:Doing devision iteration 0
DEBUG:root:Division : 3/4
INFO:root:Doing devision iteration 1
DEBUG:root:Division : 2/0
ERROR:root:Error, Division Failed
Traceback (most recent call last):
  File "<ipython-input-10-a904db1e3e23>", line 8, in my_division
    result = dividend / divisor
ZeroDivisionError: division by zero
INFO:root:Doing devision iteration 2
DEBUG:root:Division : 3/5
INFO:root:Doing devision iteration 3
DEBUG:root:Division : 10/1
ERROR:root:Error, Division Failed
Traceback (most recent call last):
  File "<ipython-input-10-a904db1e3e23>", line 8, in my_division
    result = dividend / divisor
TypeError: unsupported operand type(s) for /: 'str' and 'int'
[0.75, None, 0.6, None]
```

What is different with logging?

- We have more structure, and easier parsing
- The logging module provides some extra information (Logger, Level, and Formatting)
- We handle exception essentially for free.

Dimensions of a logging message



How does the logging module represent this configuration

Logging Aspects	Question	Description
Severity	How Important?	Log Level: DEBUG, INFO, WARN, ERROR, CRITICAL
Origin/Logger	Where?	Logger Name: ROOT, module name, etc...
Context	When?	Technical attributes, set by the logging library (time, module,...) and manual defined context (session, ...)
Message	What?	Combination of text, attributes and context
Format	How does it look like?	Template to build message
Select / Route	Who should know?	What gets send to which handler
Destination	How do we store/send/show it?	To what file, mail, socket, aggregator are the messages sent?

Back to code. How does logging work?

```
import log1;logging = log1.get_clean_logging() # this would be import logging outside  
this notebook  
  
logging.debug("Find me in the log")  
logging.info("I am hidden")  
logging.warn("I am here")  
logging.error("As am I")  
try:  
    1/0;  
except:  
    logging.exception(" And I")  
logging.critical("Me, of course")
```

```
WARNING:root:I am here  
ERROR:root:As am I  
ERROR:root: And I  
Traceback (most recent call last):  
  File "<ipython-input-12-75f8227eec02>", line 8, in <module>  
    1/0;  
ZeroDivisionError: division by zero  
CRITICAL:root:Me, of course
```

More complex Logging setup with basicConfig

```
import log1;logging = log1.get_clean_logging()

datefmt = "%Y-%m-%d %H:%M:%S"
msgfmt = "%(asctime)s,%(msecs)03d %(levelname)-10s %(name)-15s : %(message)s"
logging.basicConfig(level=logging.DEBUG, format=msgfmt, datefmt=datefmt)
logging.debug("Now I show up ")
logging.info("Now this is %s logging!","good")
logging.warn("I am here. %-4i + %-4i = %i",1,3,1+3)
logging.error("As am I")
try:
    1/0;
except:
    logging.exception(" And I")
```

```
2015-07-19 20:19:55,551 DEBUG      root      : Now I show up
2015-07-19 20:19:55,552 INFO       root      : Now this is good logging!
2015-07-19 20:19:55,552 WARNING    root      : I am here. 1      + 3      = 4
2015-07-19 20:19:55,552 ERROR      root      : As am I
2015-07-19 20:19:55,553 ERROR      root      : And I
Traceback (most recent call last):
  File "<ipython-input-13-63765f2f7e9f>", line 12, in <module>
    1/0;
ZeroDivisionError: division by zero
```

Some (personal) remarks about basicConfig()

- basicConfig does save you some typing, but I would go for the 'normal' setup.
 - it is a matter of personal taste.
- The normal setup makes the structure clearer.
- Keep in mind that basicConfig() is meant to be called once...

Using the standard configuration

slightly longer, same results....

```
import log1, json, logging.config;logging = log1.get_clean_logging()
datefmt = "%Y-%m-%d %H:%M:%S"
msgfmt = "%(asctime)s,%(msecs)03d %(levelname)-6s %(name)-10s : %(message)s"

log = logging.getLogger()
log.setLevel(logging.DEBUG)
lh = logging.StreamHandler()
lf = logging.Formatter(fmt=msgfmt, datefmt=datefmt)
lh.setFormatter(lf)
log.addHandler(lh)

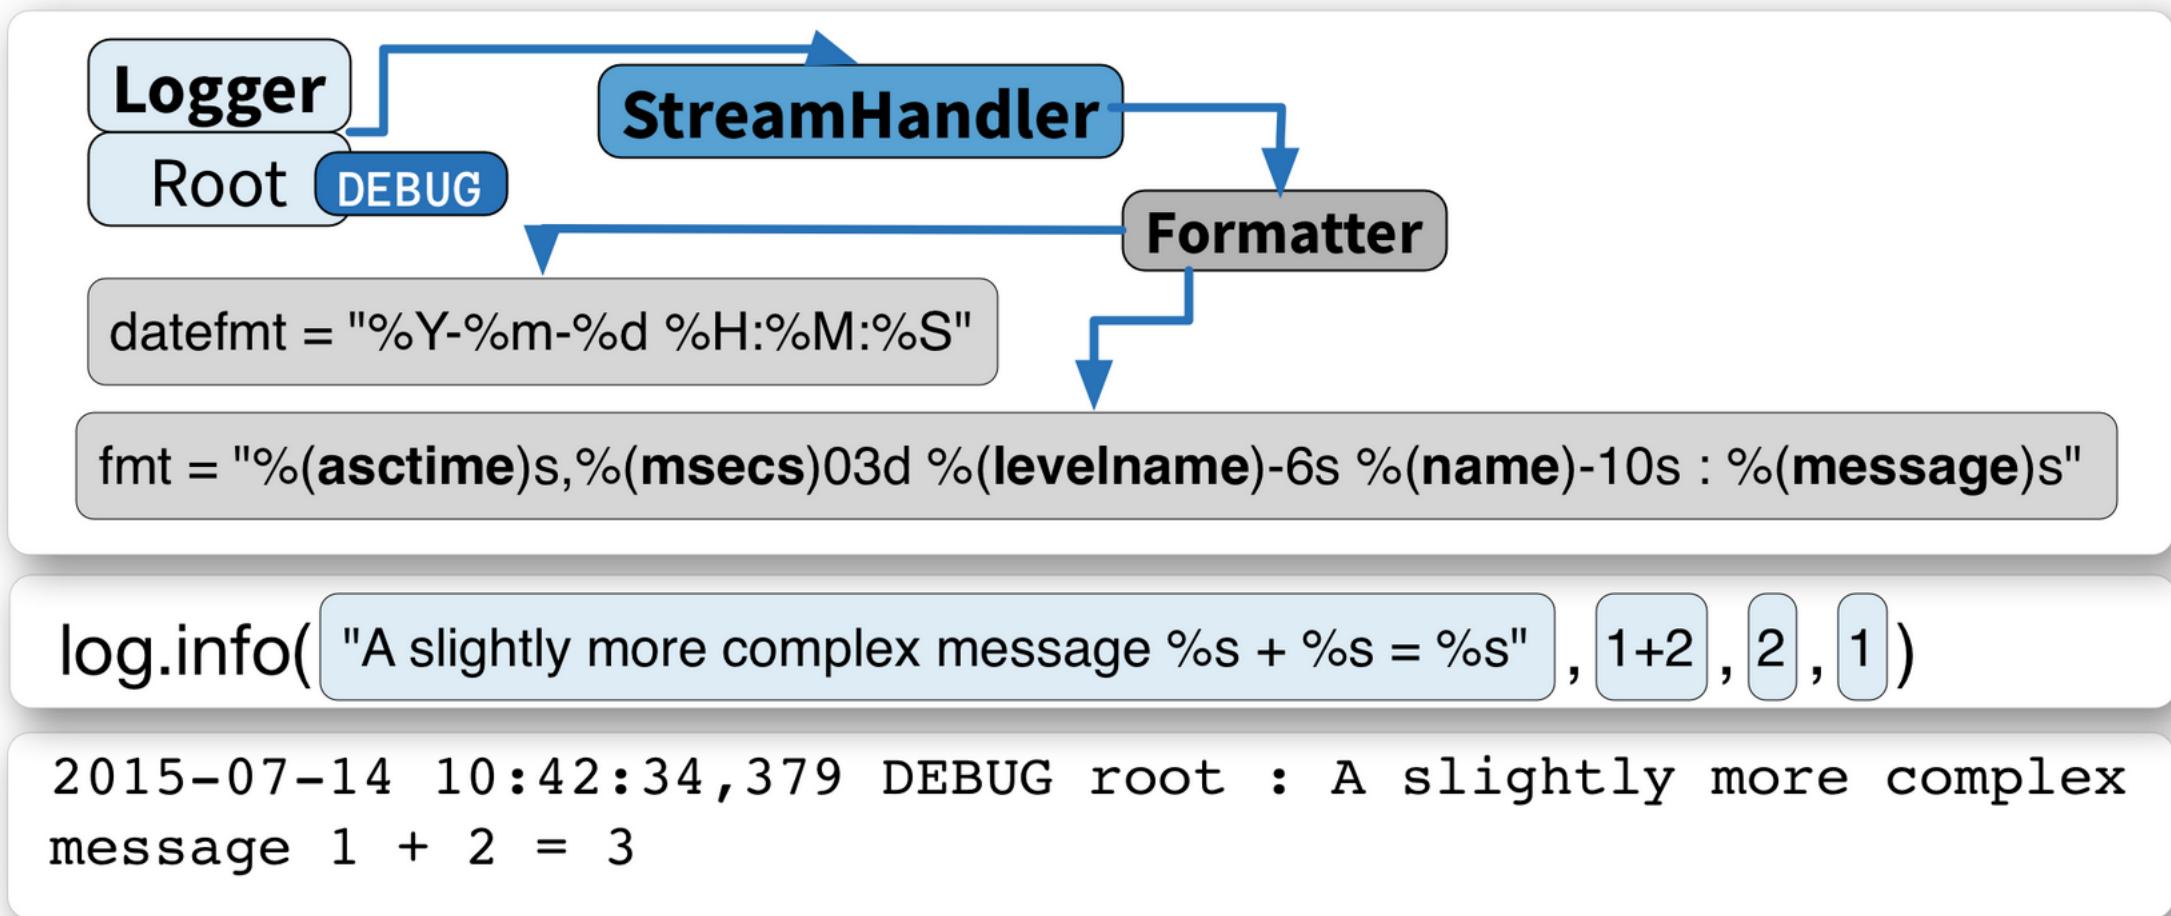
log.info("Now this is %s logging!", "good")
log.debug("A slightly more complex message %s + %s = %s",1,2,1+2)
```

```
2015-07-19 20:19:55,571 INFO    root      : Now this is good logging!
2015-07-19 20:19:55,572 DEBUG    root      : A slightly more complex message 1 + 2
= 3
```

Now, back the some theory. What have we build?



How do we get from the configuration to the log message?



Formatting : Attributes available for the logging call

Attribute	Description
args	Tuple of arguments passed to the logging call
asctime	Log record creation time, formatted
created	Log record creation time, seconds since the Epoch
exc_info	Exception information / stack trace, if any
filename	Filename portion of pathname for the logging module
funcName	Name of function containing the logging call
levelname	Name of Logging Level
levelno	Number of Logging Level
lineno	Line number in source code for the logging call
module	Module (name portion of filename).
message	Logged message
name	Name of the logger used to log the call.
pathname	pathname of source file
process	Process ID
processName	Process name
thread	Thread ID
threadName	Thread name

Using dictConfig()

```
import log1, json, logging.config;logging = log1.get_clean_logging()
conf_dict = {
    'version': 1,
    'disable_existing_loggers': True,
    'formatters': {
        'longformat': {
            'format': "%(asctime)s,%(msecs)03d %(levelname)-10s %(name)-15s : %(message)s",
            'datefmt': "%Y-%m-%d %H:%M:%S"},},
    'handlers': {
        'console': {
            'class': 'logging.StreamHandler',
            'formatter': "longformat"},},
    'loggers':{
        '': {
            'level': 'DEBUG',
            'handlers': ['console']}}
logging.config.dictConfig(conf_dict)
log = logging.getLogger()
log.info("Now this is %s logging!", "good")
```

2015-07-19 20:19:55,602 INFO

root

: Now this is good logging!

Adding a Filehandler to the Logger

```
import log1, json, logging.config;logging = log1.get_clean_logging()
base_config = json.load(open("conf_dict.json"))

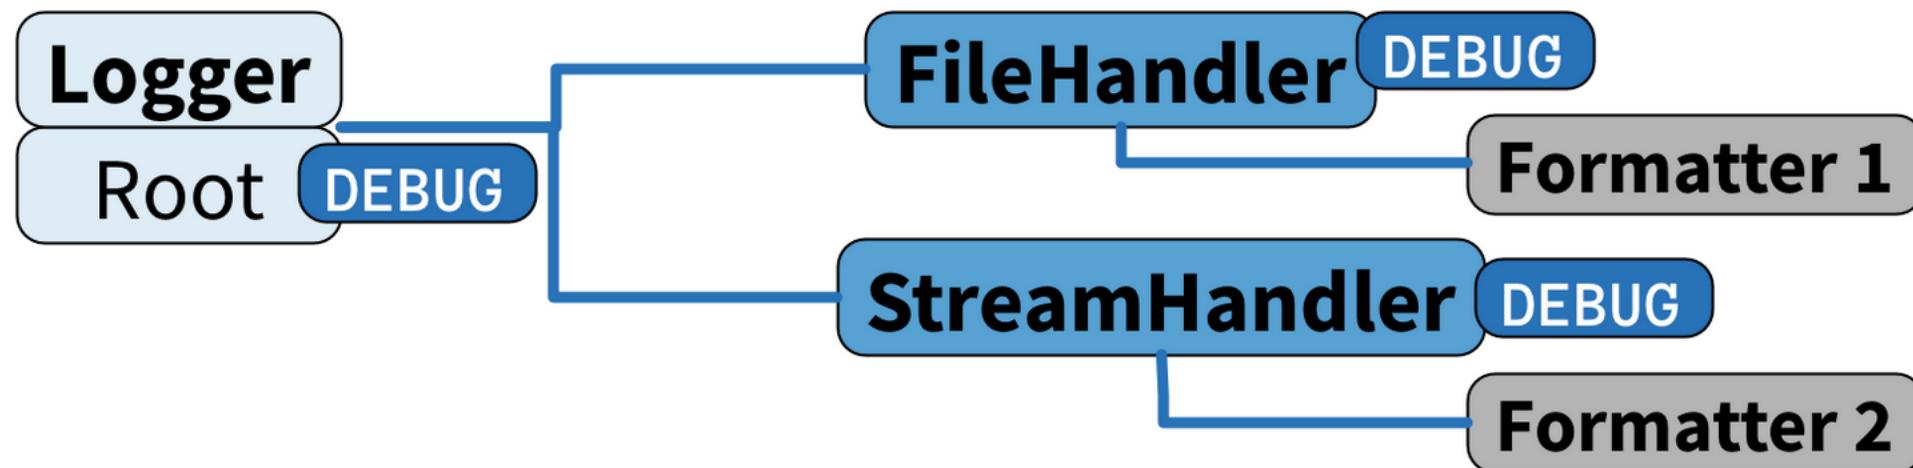
base_config['handlers']['logfile'] = {
    'class' : 'logging.FileHandler',
    'mode' : 'w',
    'filename' : 'logfile.txt',
    'formatter': "longformat"}
base_config['loggers']['']['handlers'].append('logfile')
logging.config.dictConfig(base_config)
log = logging.getLogger()
log.info("Now this is %s logging!", "good")
!cat logfile.txt
```

```
2015-07-19 20:19:55,618 INFO      root      : Now this is good logging!
```

```
2015-07-19 20:19:55,618 INFO      root      : Now this is good logging!
```

Another look at the logging object tree

now with more FileHandler



Set the Level on the FileHandler

```
import log1, json, logging.config;logging = log1.get_clean_logging()

file_config = json.load(open("conf_dict_with_file.json"))
file_config['handlers'][ 'logfile' ][ 'level' ] = "WARN"
logging.config.dictConfig(file_config)
log = logging.getLogger()
log.info("Now this is %s logging!", "good")
log.warning("Now this is %s logging!", "worrisome")
!cat logfile.txt
```

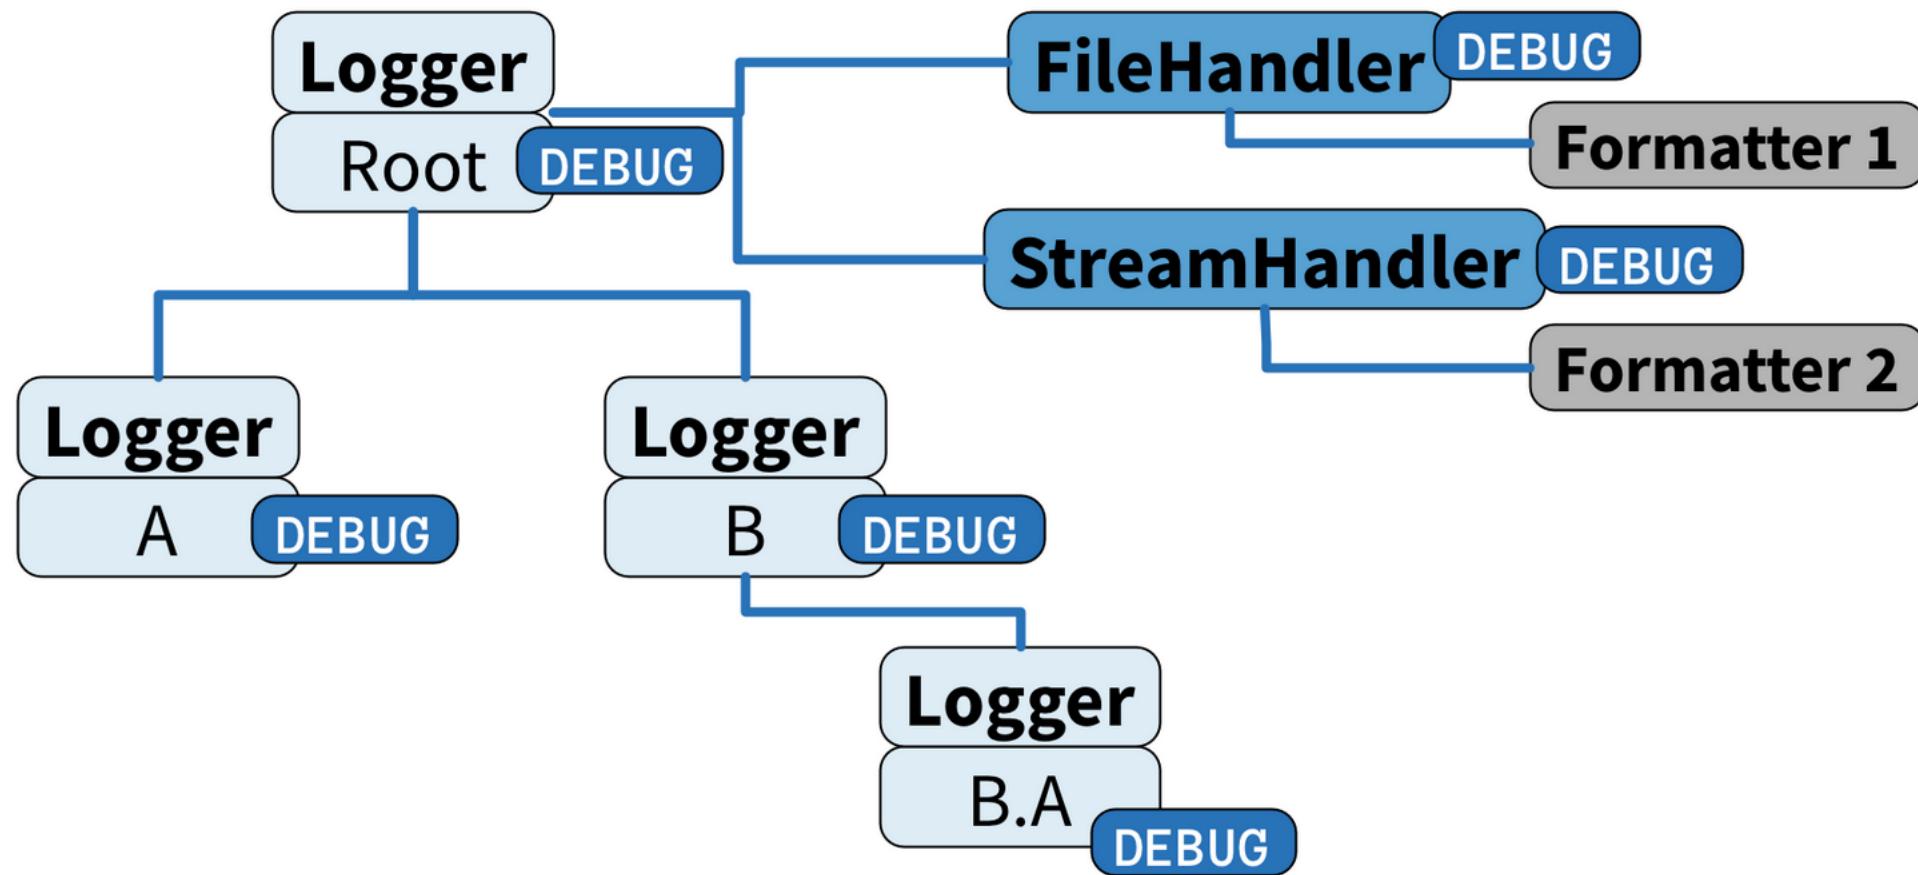
2015-07-19 20:19:55,744 INFO	root	: Now this is good logging!
2015-07-19 20:19:55,745 WARNING	root	: Now this is worrisome logging!
2015-07-19 20:19:55,745 WARNING	root	: Now this is worrisome logging!

Adding Child Loggers under the Root

```
import log1,json,logging.config;logging = log1.get_clean_logging()
logging.config.dictConfig(json.load(open("conf_dict.json")))
log = logging.getLogger("")
child_A = logging.getLogger("A")
child_B = logging.getLogger("B")
child_B_A = logging.getLogger("B.A")
log.info("Now this is %s logging!", "good")
child_A.info("Now this is more logging!")
log.warning("Now this is %s logging!", "worrisome")
```

2015-07-19 20:19:55,865	INFO	root	: Now this is good logging!
2015-07-19 20:19:55,866	INFO	A	: Now this is more logging!
2015-07-19 20:19:55,867	WARNING	root	: Now this is worrisome logging!

Looking at the tree of logging objects



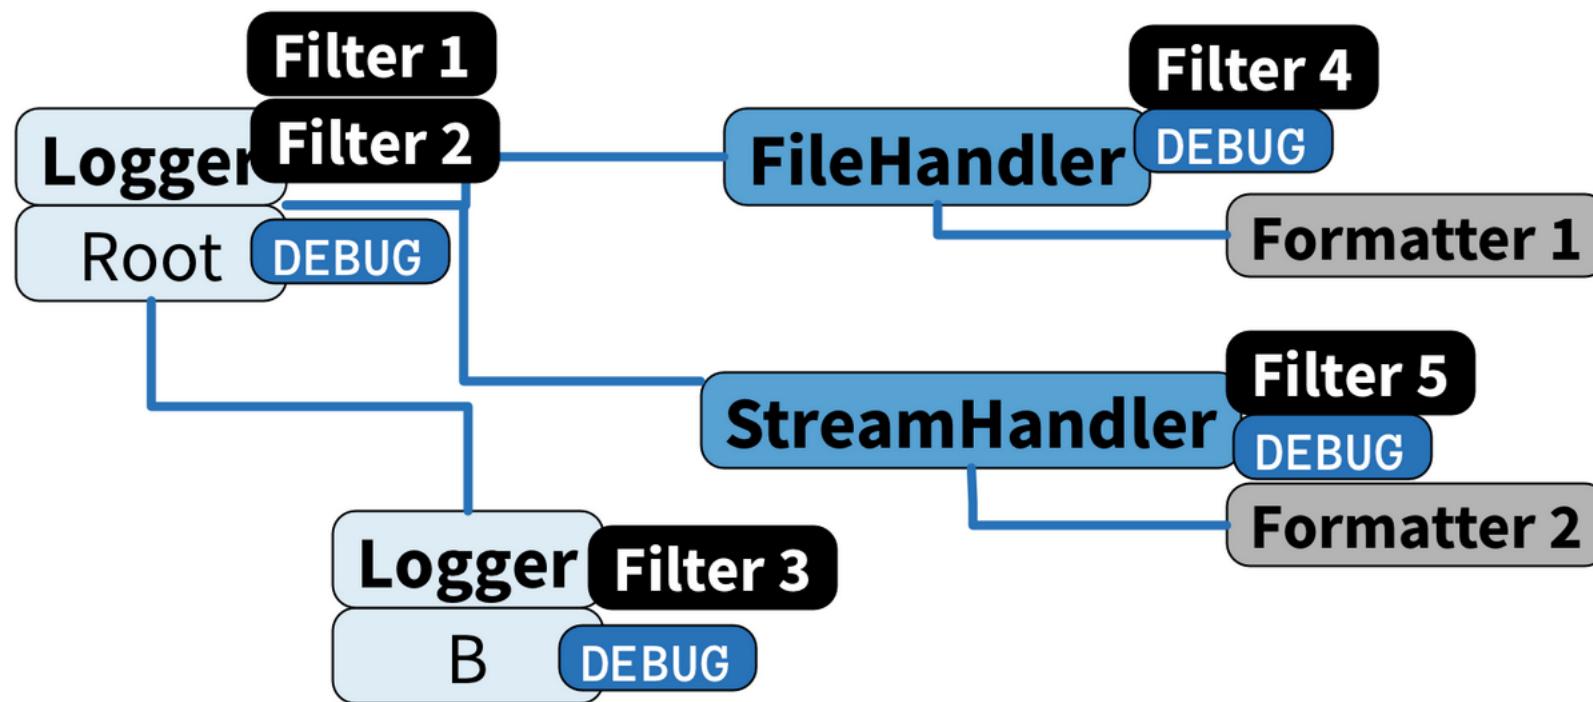
Best Practices for the Logging Tree

- Use `.getLogger(__name__)` per module to define loggers under the root logger
- Set propagate to True on each Logger
- Attach Handlers and Filters as needed to control output from the Logging hierarchy

Filter - Now that things are getting complicated

- With more loggers and handlers in the tree of logging objects, things are getting complicated
- We may not want every logger to send log records to every filter
 - The logging level gives us some control, there are limits
- Filters are one solution to this problem
- Filter can also add information to records, thus helping with structured logging

Using Filters



Filter = object with a filter(record) method.

Returns True or False.

Can modify in place to add / change/ remove information

An Example for using Filter

```
import log1,json,logging.config;logging = log1.get_clean_logging()
logging.config.dictConfig(json.load(open("conf_dict.json")))

def log_filter(rec): # Callables work with 3.2 and later
    if 'please' in rec.msg.lower():
        return True
    return False

log = logging.getLogger("")
log.addFilter(log_filter)
child_A = logging.getLogger("A")

log.info("Just log me")
child_A.info("Just log me")
log.info("Hallo, Please log me")
```

2015-07-20 08:01:55,108 INFO

A

: Just log me

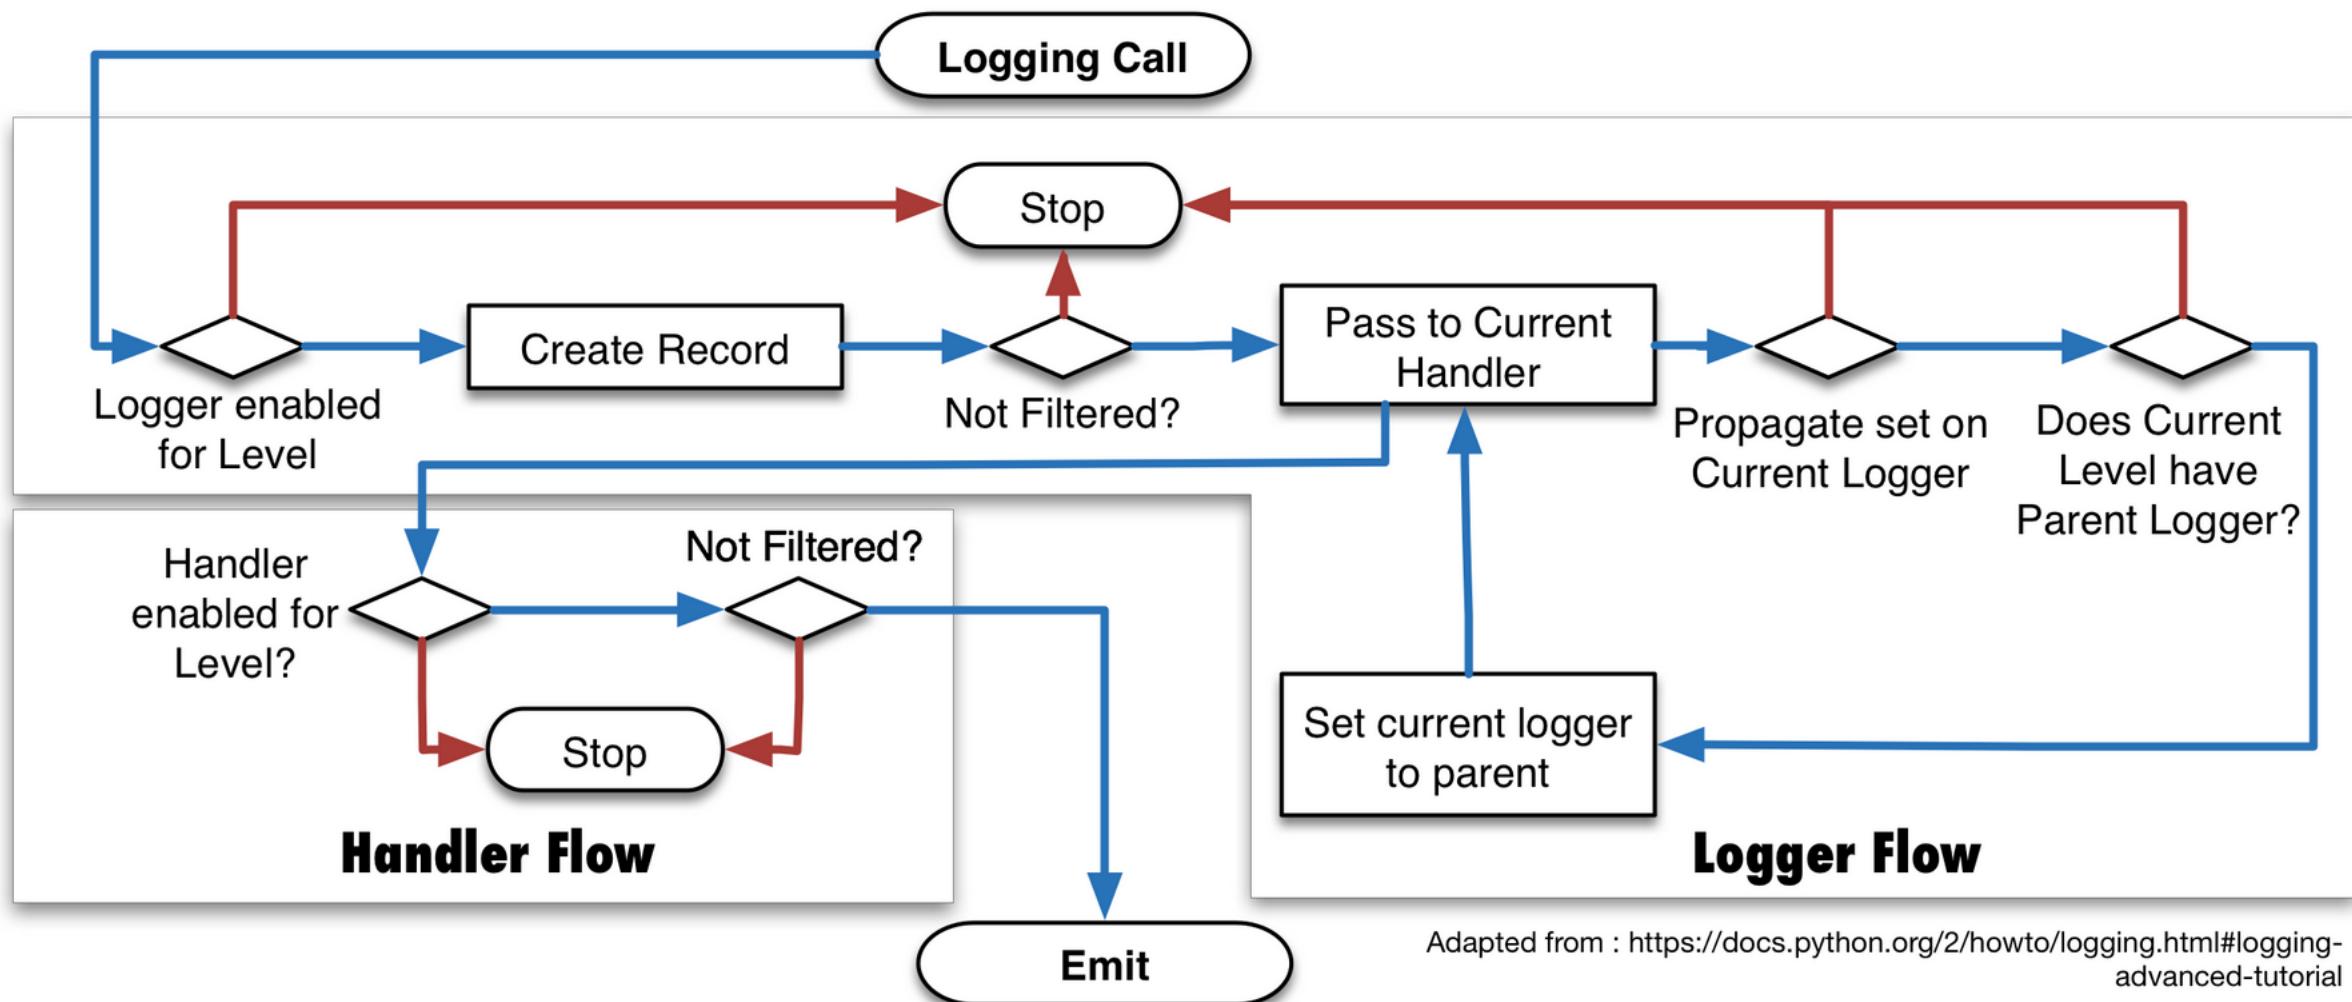
2015-07-20 08:01:55,108 INFO

root

: Hallo, Please log me

The Way of a logging record

Or: Why does Logger A's "Now this is more logging" get logged



A second Examples for Filters, in the Handler

```
import log1, json, logging.config;logging = log1.get_clean_logging()
datefmt = "%Y-%m-%d %H:%M:%S"
msgfmt = "%(asctime)s,%(msecs)03d %(levelname)-6s %(name)-10s : %(message)s"
log_reg = None
def handler_filter(rec): # Callables work with 3.2 and later
    global log_reg
    if 'please' in rec.msg.lower():
        rec.msg = rec.msg + " (I am nice)" # Changing the record
        rec.args = (rec.args[0].upper(), rec.args[1] + 10)
        rec.__dict__['custom_name'] = "Important context information"
        log_reg = rec
    return True
return False
log = logging.getLogger()
lh = logging.StreamHandler()
lf = logging.Formatter(fmt=msgfmt, datefmt=datefmt)
lh.setFormatter(lf)
log.addHandler(lh)
lh.addFilter(handler_filter)
log.warn("I am a bold Logger","good")
log.warn("Hi, I am %s. I am %i seconds old. Please log me","Loggy", 1)
```

2015-07-19 20:19:55,905 WARNING root : Hi, I am LOGGY. I am 11 seconds old.
Please log me (I am nice)

Things you might want to know (if we still have some time)

A short look at our LogRecord

```
print(log_reg)
log_reg.__dict__
```

```
<LogRecord: root, 30, <ipython-input-20-d1d101ab918f>, 25, "Hi, I am %s. I am %i
seconds old. Please log me (I am nice)">
```

```
{'args': ('LOGGY', 11),
 'asctime': '2015-07-19 20:19:55',
 'created': 1437329995.905689,
 'custom_name': 'Important context information',
 'exc_info': None,
 'exc_text': None,
 'filename': '<ipython-input-20-d1d101ab918f>',
 'funcName': '<module>',
 'levelname': 'WARNING',
 'levelno': 30,
 'lineno': 25,
 'message': 'Hi, I am LOGGY. I am 11 seconds old. Please log me (I am nice)',
 'module': '<ipython-input-20-d1d101ab918f>',
 'msecs': 905.689001083374,
 'msg': 'Hi, I am %s. I am %i seconds old. Please log me (I am nice)',
 'name': 'root',
 'pathname': '<ipython-input-20-d1d101ab918f>',
 'process': 1644,
 'processName': 'MainProcess',
 'relativeCreated': 1.280069351196289,
 'stack_info': None,
 'thread': 140735243608832,
```

Logging Performance - Slow, but fast enough

Scenario (10000 Call, 3 Logs per call)	Runtime
Full Logging with buffered writes	3.096s
Disable Caller information	2.868s
Check Logging Lvl before Call, Logging disabled	0.186s
Logging module level disabled	0.181s
No Logging calls at all	0.157s

Getting the current Logging Tree

```
config = json.load(open("conf_dict_with_file.json"))
logging.config.dictConfig(config)
import requests
import logging_tree
logging_tree.printout()
```

```
<--"""
    Level DEBUG
    Handler <logging.StreamHandler object at 0x107542780>
        Formatter fmt='%(asctime)s,%(msecs)03d %(levelname)-10s %(name)-15s : %(message)s' datefmt='%Y-%m-%d %H:%M:%S'
    Handler <logging.FileHandler object at 0x107542ef0>
        Formatter fmt='%(asctime)s,%(msecs)03d %(levelname)-10s %(name)-15s : %(message)s' datefmt='%Y-%m-%d %H:%M:%S'
    |
o<--"requests"
    Level DEBUG
    Disabled
    Handler <logging.FileHandler object at 0x10759f9e8>
```

Reconfiguration

- It is possible to change the logging configuration at runtime
- It is even part of the standard library
- Still, some caution is in order

Reloading the configuration *can* disable the existing loggers

```
import log1,json,logging,logging.config;logging = log1.get_clean_logging()

#Load Config, define a child logger (could also be a module)
logging.config.dictConfig(json.load(open("conf_dict_with_file.json")))
child_log = logging.getLogger("somewhere")

#Reload Config
logging.config.dictConfig(json.load(open("conf_dict_with_file.json")))

#Our childlogger was disabled
child_log.info("Now this is %s logging!", "good")
```

Reloading can happen in place

```
import log1, json, logging, logging.config;logging = log1.get_clean_logging()

config = json.load(open("conf_dict_with_file.json"))
#Load Config, define a child logger (could also be a module)

logging.config.dictConfig(config)
child_log = logging.getLogger("somewhere")
config['disable_existing_loggers'] = False
#Reload Config
logging.config.dictConfig(config)

#Our childlogger was disabled
child_log.info("Now this is %s logging!", "good")
```

2015-07-19 20:20:42,290 INFO somewhere : Now this is good logging!